

fetal-health-over-sampled-1

August 10, 2024

```
[34]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
from imblearn.over_sampling import SMOTE
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
```

```
[35]: data=pd.read_csv('/kaggle/input/fetal-health-classification/fetal_health.csv')
data.head()
```

```
[35]: baseline value accelerations fetal_movement uterine_contractions \
0          120.0          0.000          0.0          0.000
1          132.0          0.006          0.0          0.006
2          133.0          0.003          0.0          0.008
3          134.0          0.003          0.0          0.008
4          132.0          0.007          0.0          0.008

light_decelerations severe_decelerations prolonged_decelerations \
0          0.000          0.0          0.0
1          0.003          0.0          0.0
2          0.003          0.0          0.0
3          0.003          0.0          0.0
4          0.000          0.0          0.0

abnormal_short_term_variability mean_value_of_short_term_variability \
0          73.0          0.5
1          17.0          2.1
2          16.0          2.1
3          16.0          2.4
4          16.0          2.4

percentage_of_time_with_abnormal_long_term_variability ... histogram_min \
0          43.0          ...          62.0
1          0.0          ...          68.0
2          0.0          ...          68.0
```

3	0.0	...	53.0
4	0.0	...	53.0

	histogram_max	histogram_number_of_peaks	histogram_number_of_zeroes	\
0	126.0	2.0	0.0	
1	198.0	6.0	1.0	
2	198.0	5.0	1.0	
3	170.0	11.0	0.0	
4	170.0	9.0	0.0	

	histogram_mode	histogram_mean	histogram_median	histogram_variance	\
0	120.0	137.0	121.0	73.0	
1	141.0	136.0	140.0	12.0	
2	141.0	135.0	138.0	13.0	
3	137.0	134.0	137.0	13.0	
4	137.0	136.0	138.0	11.0	

	histogram_tendency	fetal_health
0	1.0	2.0
1	0.0	1.0
2	0.0	1.0
3	1.0	1.0
4	1.0	1.0

[5 rows x 22 columns]

```
[36]: print("Rows:",data.shape[0])
      print("Columns:",data.shape[1])
```

Rows: 2126
Columns: 22

```
[37]: data.describe()
```

```
[37]:
```

	baseline value	accelerations	fetal_movement	uterine_contractions	\
count	2126.000000	2126.000000	2126.000000	2126.000000	
mean	133.303857	0.003178	0.009481	0.004366	
std	9.840844	0.003866	0.046666	0.002946	
min	106.000000	0.000000	0.000000	0.000000	
25%	126.000000	0.000000	0.000000	0.002000	
50%	133.000000	0.002000	0.000000	0.004000	
75%	140.000000	0.006000	0.003000	0.007000	
max	160.000000	0.019000	0.481000	0.015000	

	light_decelerations	severe_decelerations	prolongued_decelerations	\
count	2126.000000	2126.000000	2126.000000	
mean	0.001889	0.000003	0.000159	

std	0.002960	0.000057	0.000590
min	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000
75%	0.003000	0.000000	0.000000
max	0.015000	0.001000	0.005000

	abnormal_short_term_variability	mean_value_of_short_term_variability \
count	2126.000000	2126.000000
mean	46.990122	1.332785
std	17.192814	0.883241
min	12.000000	0.200000
25%	32.000000	0.700000
50%	49.000000	1.200000
75%	61.000000	1.700000
max	87.000000	7.000000

	percentage_of_time_with_abnormal_long_term_variability ... \
count	2126.00000 ...
mean	9.84666 ...
std	18.39688 ...
min	0.00000 ...
25%	0.00000 ...
50%	0.00000 ...
75%	11.00000 ...
max	91.00000 ...

	histogram_min	histogram_max	histogram_number_of_peaks \
count	2126.000000	2126.000000	2126.000000
mean	93.579492	164.025400	4.068203
std	29.560212	17.944183	2.949386
min	50.000000	122.000000	0.000000
25%	67.000000	152.000000	2.000000
50%	93.000000	162.000000	3.000000
75%	120.000000	174.000000	6.000000
max	159.000000	238.000000	18.000000

	histogram_number_of_zeroes	histogram_mode	histogram_mean \
count	2126.000000	2126.000000	2126.000000
mean	0.323612	137.452023	134.610536
std	0.706059	16.381289	15.593596
min	0.000000	60.000000	73.000000
25%	0.000000	129.000000	125.000000
50%	0.000000	139.000000	136.000000
75%	0.000000	148.000000	145.000000
max	10.000000	187.000000	182.000000

	histogram_median	histogram_variance	histogram_tendency	fetal_health
count	2126.000000	2126.000000	2126.000000	2126.000000
mean	138.090310	18.808090	0.320320	1.304327
std	14.466589	28.977636	0.610829	0.614377
min	77.000000	0.000000	-1.000000	1.000000
25%	129.000000	2.000000	0.000000	1.000000
50%	139.000000	7.000000	0.000000	1.000000
75%	148.000000	24.000000	1.000000	1.000000
max	186.000000	269.000000	1.000000	3.000000

[8 rows x 22 columns]

[38]: data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2126 entries, 0 to 2125
Data columns (total 22 columns):
#   Column                                     Non-Null Count
Dtype
---  ---
0    baseline value                           2126 non-null
float64
1    accelerations                           2126 non-null
float64
2    fetal_movement                           2126 non-null
float64
3    uterine_contractions                     2126 non-null
float64
4    light_decelerations                      2126 non-null
float64
5    severe_decelerations                     2126 non-null
float64
6    prolonged_decelerations                  2126 non-null
float64
7    abnormal_short_term_variability           2126 non-null
float64
8    mean_value_of_short_term_variability      2126 non-null
float64
9    percentage_of_time_with_abnormal_long_term_variability 2126 non-null
float64
10   mean_value_of_long_term_variability       2126 non-null
float64
11   histogram_width                           2126 non-null
float64
12   histogram_min                             2126 non-null
float64
```

```

    13 histogram_max                2126 non-null
float64
    14 histogram_number_of_peaks    2126 non-null
float64
    15 histogram_number_of_zeroes    2126 non-null
float64
    16 histogram_mode               2126 non-null
float64
    17 histogram_mean               2126 non-null
float64
    18 histogram_median              2126 non-null
float64
    19 histogram_variance            2126 non-null
float64
    20 histogram_tendency            2126 non-null
float64
    21 fetal_health                  2126 non-null
float64
dtypes: float64(22)
memory usage: 365.5 KB

```

```
[39]: data.duplicated()
```

```

[39]: 0      False
      1      False
      2      False
      3      False
      4      False
      ...
     2121    False
     2122    False
     2123    False
     2124    False
     2125    False
      Length: 2126, dtype: bool

```

```
[40]: data.notnull()
```

```

[40]:      baseline value  accelerations  fetal_movement  uterine_contractions  \
0           True        True        True        True
1           True        True        True        True
2           True        True        True        True
3           True        True        True        True
4           True        True        True        True
...          ...          ...          ...          ...
     2121    True        True        True        True
     2122    True        True        True        True

```

2123	True	True	True	True
2124	True	True	True	True
2125	True	True	True	True

	light_decelerations	severe_decelerations	prolongued_decelerations	\
0	True	True	True	True
1	True	True	True	True
2	True	True	True	True
3	True	True	True	True
4	True	True	True	True
...
2121	True	True	True	True
2122	True	True	True	True
2123	True	True	True	True
2124	True	True	True	True
2125	True	True	True	True

	abnormal_short_term_variability	mean_value_of_short_term_variability	\
0	True	True	True
1	True	True	True
2	True	True	True
3	True	True	True
4	True	True	True
...
2121	True	True	True
2122	True	True	True
2123	True	True	True
2124	True	True	True
2125	True	True	True

	percentage_of_time_with_abnormal_long_term_variability	...	\
0	True	...	True
1	True	...	True
2	True	...	True
3	True	...	True
4	True	...	True
...
2121	True	...	True
2122	True	...	True
2123	True	...	True
2124	True	...	True
2125	True	...	True

	histogram_min	histogram_max	histogram_number_of_peaks	\
0	True	True	True	True
1	True	True	True	True
2	True	True	True	True

3	True	True	True
4	True	True	True
...
2121	True	True	True
2122	True	True	True
2123	True	True	True
2124	True	True	True
2125	True	True	True

	histogram_number_of_zeroes	histogram_mode	histogram_mean	\
0	True	True	True	
1	True	True	True	
2	True	True	True	
3	True	True	True	
4	True	True	True	
...	
2121	True	True	True	
2122	True	True	True	
2123	True	True	True	
2124	True	True	True	
2125	True	True	True	

	histogram_median	histogram_variance	histogram_tendency	fetal_health
0	True	True	True	True
1	True	True	True	True
2	True	True	True	True
3	True	True	True	True
4	True	True	True	True
...
2121	True	True	True	True
2122	True	True	True	True
2123	True	True	True	True
2124	True	True	True	True
2125	True	True	True	True

[2126 rows x 22 columns]

```
[41]: data.keys()
```

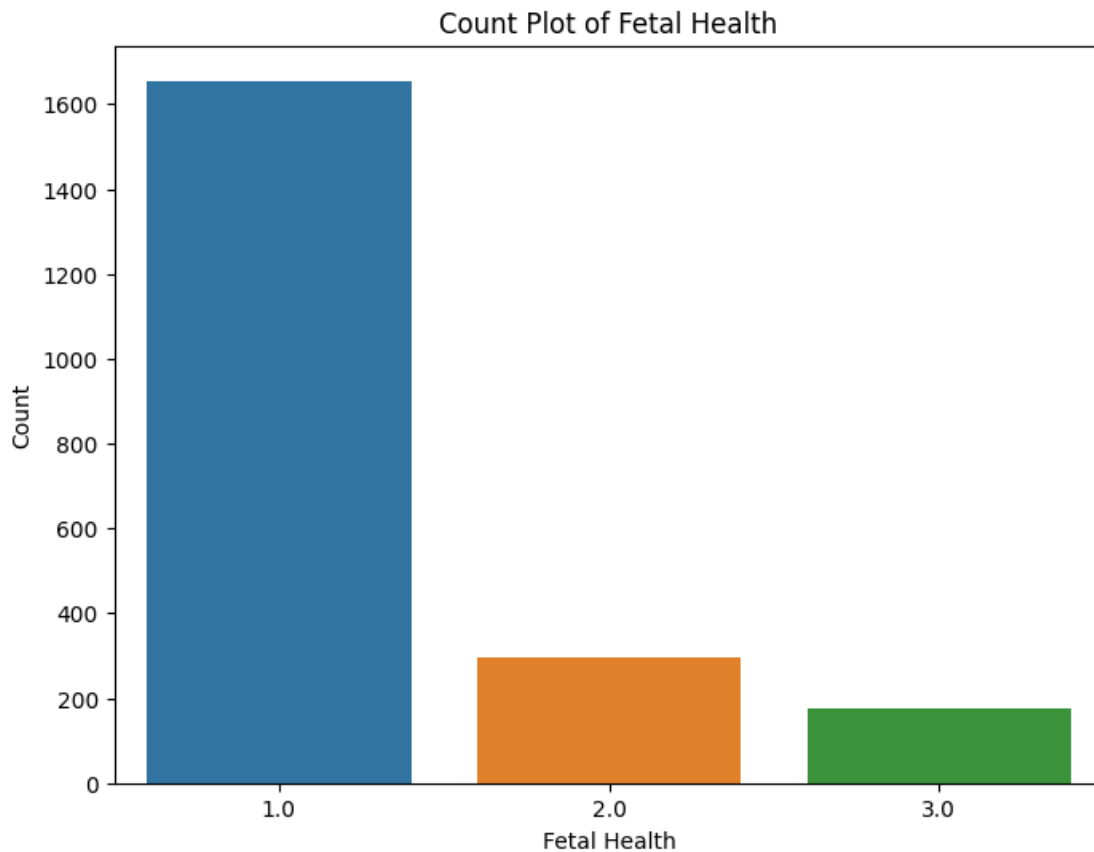
```
[41]: Index(['baseline value', 'accelerations', 'fetal_movement',
'uterine_contractions', 'light_decelerations', 'severe_decelerations',
'prolongued_decelerations', 'abnormal_short_term_variability',
'mean_value_of_short_term_variability',
'percentage_of_time_with_abnormal_long_term_variability',
'mean_value_of_long_term_variability', 'histogram_width',
'histogram_min', 'histogram_max', 'histogram_number_of_peaks',
'histogram_number_of_zeroes', 'histogram_mode', 'histogram_mean',
```

```
'histogram_median', 'histogram_variance', 'histogram_tendency',  
'fetal_health'],  
dtype='object')
```

```
[42]: # Analysing the target column  
data['fetal_health'].unique()
```

```
[42]: array([2., 1., 3.])
```

```
[43]: # Plot count plot  
plt.figure(figsize=(8, 6))  
sns.countplot(x='fetal_health', data=data)  
plt.title('Count Plot of Fetal Health')  
plt.xlabel('Fetal Health')  
plt.ylabel('Count')  
plt.show()
```



```
[44]: import matplotlib.pyplot as plt  
import numpy as np
```



```

# Define the figure size
plt.figure(figsize=(15, 8)) # Adjust the width and height according to your
    ↳ preference

# List of features
features = ['baseline value', 'accelerations', 'fetal_movement',
            'uterine_contractions', 'light_decelerations',
    ↳ 'severe_decelerations',
            'prolongued_decelerations', 'abnormal_short_term_variability',
            'mean_value_of_short_term_variability',
            'percentage_of_time_with_abnormal_long_term_variability',
            'mean_value_of_long_term_variability', 'histogram_width',
            'histogram_min', 'histogram_max', 'histogram_number_of_peaks',
            'histogram_number_of_zeroes', 'histogram_mode', 'histogram_mean',
            'histogram_median', 'histogram_variance', 'histogram_tendency']

# Plotting each feature against the target variable 'fetal_health'
num_features = len(features)
num_rows = num_features // 3 + (1 if num_features % 3 != 0 else 0) # Calculate
    ↳ the number of rows needed
fig, axes = plt.subplots(nrows=num_rows, ncols=3, figsize=(15, 4*num_rows))

for i, feature in enumerate(features):
    row = i // 3 # Calculate the row index
    col = i % 3 # Calculate the column index
    ax = axes[row, col] if num_rows > 1 else axes[col] # Select the
    ↳ appropriate axis
    ax.scatter(data[feature], data['fetal_health'], cmap='viridis')
    ax.set_ylabel('Fetal Health')
    ax.set_xlabel(feature)
    ax.set_title(f'Scatter Plot of {feature} against Fetal Health')
    ax.grid(True)

# Remove empty subplots
if num_features % 3 != 0:
    for j in range(num_features % 3, 3):
        fig.delaxes(axes[num_rows - 1, j])

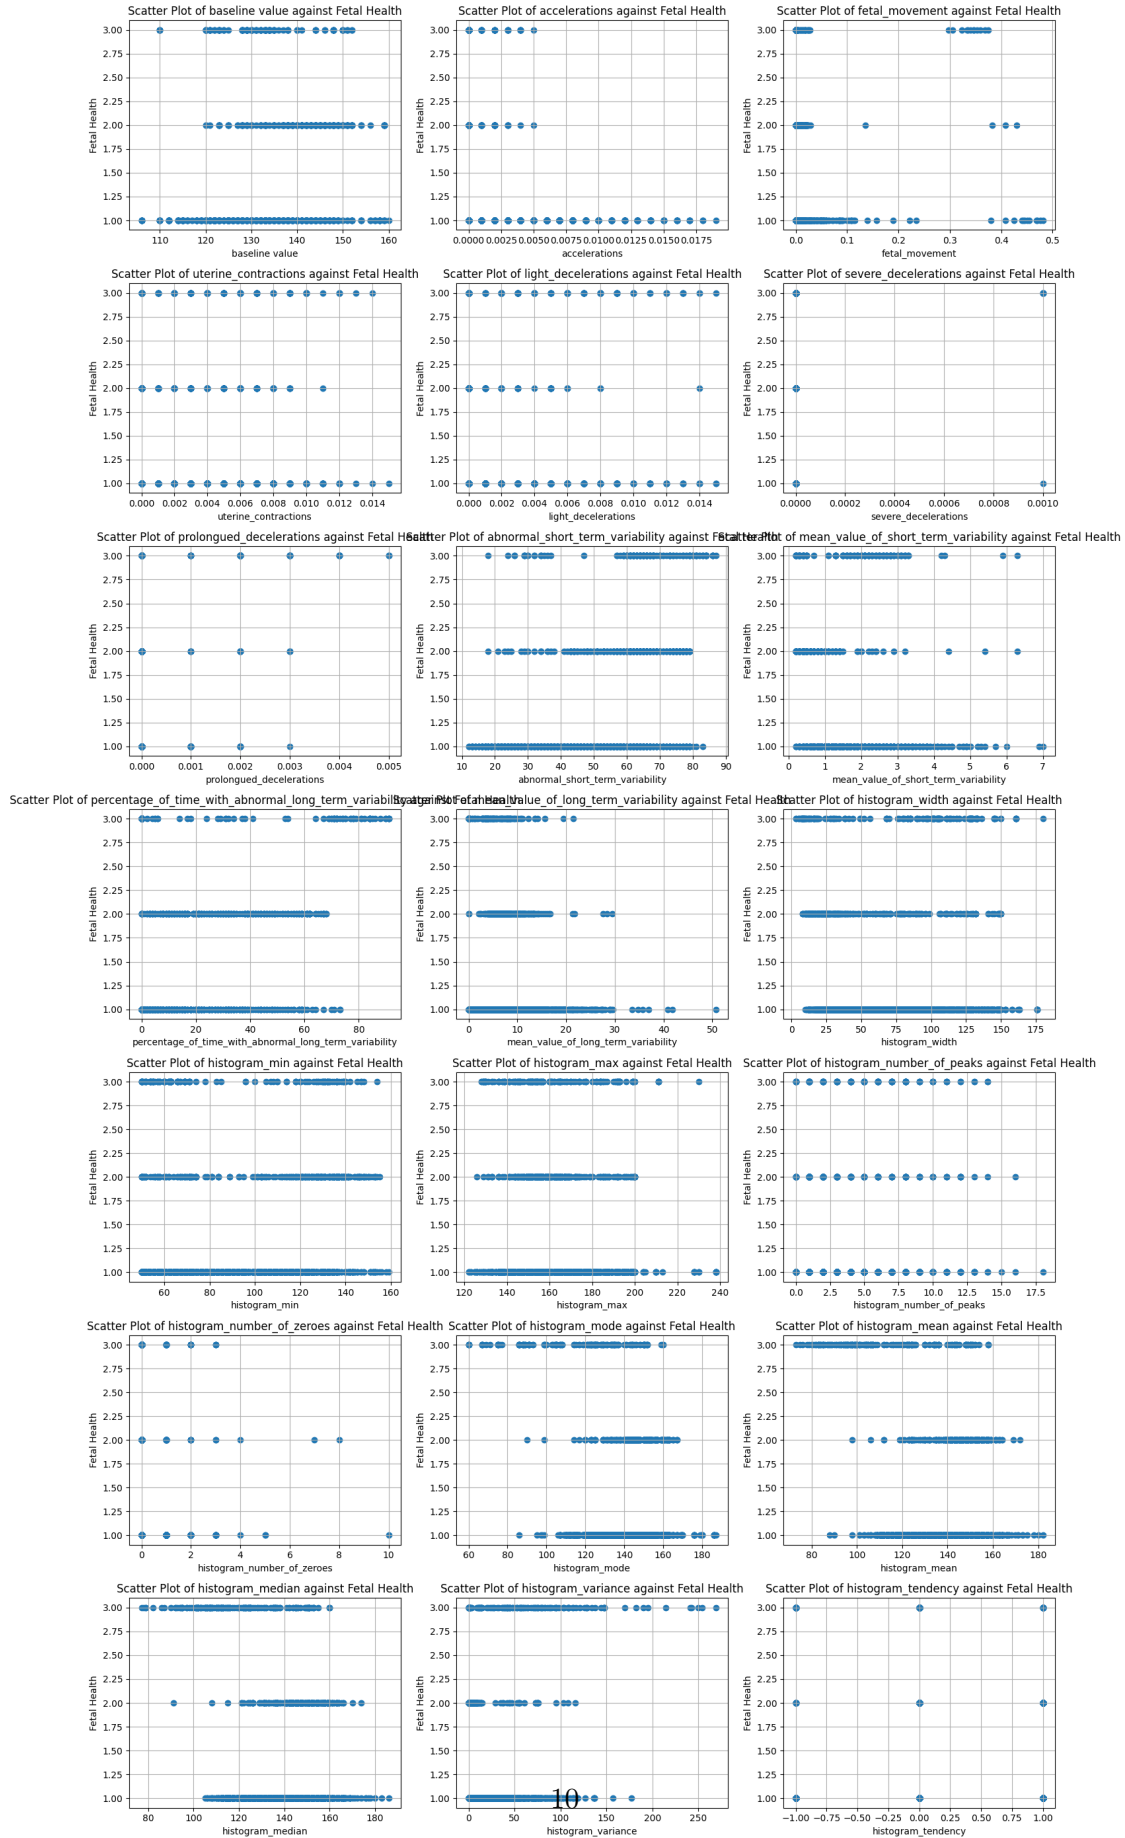
plt.tight_layout()
plt.show()

```

/tmp/ipykernel_33/2953099270.py:27: UserWarning: No data for colormapping provided via 'c'. Parameters 'cmap' will be ignored

```
ax.scatter(data[feature], data['fetal_health'], cmap='viridis')
```

<Figure size 1500x800 with 0 Axes>




```
[45]: # Assuming 'df' is your DataFrame and 'target_column' is the column containing
      ↪ class labels
class_distribution = data['fetal_health'].value_counts()
print("Class distribution:")
print(class_distribution)

# Plot class distribution
plt.figure(figsize=(8, 6))
class_distribution.plot(kind='bar', color='skyblue')
plt.title('Class Distribution')
plt.xlabel('Class')
plt.ylabel('Frequency')
plt.xticks(rotation=45)
plt.show()

# Calculate imbalance ratio
imbalance_ratio = class_distribution.min() / class_distribution.max()
print("Imbalance ratio:", imbalance_ratio)
```

```
Class distribution:
fetal_health
1.0      1655
2.0       295
3.0       176
Name: count, dtype: int64
```



Imbalance ratio: 0.10634441087613293

```
[46]: # Assuming 'df' is your DataFrame and 'fetal_health' is the target variable
X = data.drop(columns=['fetal_health']) # Features
y = data['fetal_health'] # Target variable

# Initialize SMOTE
smote = SMOTE()

# Perform SMOTE oversampling
X_resampled, y_resampled = smote.fit_resample(X, y)

# Convert the resampled data to DataFrame
df_resampled = pd.concat([pd.DataFrame(X_resampled, columns=X.columns), pd.
    DataFrame(y_resampled, columns=['fetal_health'])], axis=1)

# Check the class distribution after oversampling
print("Class distribution after oversampling:")
print(df_resampled['fetal_health'].value_counts())
```

Class distribution after oversampling:

fetal_health

2.0 1655


1.0 1655

3.0 1655

Name: count, dtype: int64

```
[47]: df_resampled
      x=df_resampled.drop('fetal_health',axis=1)
      y=df_resampled['fetal_health']
```

```
[48]: # Plot count plot
      plt.figure(figsize=(8, 6))
      sns.countplot(x='fetal_health', data=df_resampled)
      plt.title('Count Plot after Oversampling')
      plt.xlabel('Fetal Health')
      plt.ylabel('Count')
      plt.show()
```



```
[49]: features = ['baseline value', 'accelerations', 'fetal_movement',
↳ 'uterine_contractions', 'light_decelerations',
↳
↳ 'severe_decelerations', 'prolongued_decelerations', 'abnormal_short_term_variability', 'mean_v
↳
↳ 'percentage_of_time_with_abnormal_long_term_variability', 'mean_value_of_long_term_variabili
↳
↳ 'histogram_min', 'histogram_max', 'histogram_number_of_peaks', 'histogram_number_of_zeroes', 'h
    'histogram_median', 'histogram_variance', 'histogram_tendency']

# Separating out the features
x = df_resampled.loc[:, features].values

# Separating out the target
y = df_resampled.loc[:, ['fetal_health']].values
```

```
[50]: x = StandardScaler().fit_transform(df_resampled)
```

```
[51]: y.shape
```

```
[51]: (4965, 1)
```

```
[52]: from sklearn.decomposition import PCA

pca = PCA(n_components=5)

principalComponents = pca.fit_transform(x)

principalDf = pd.DataFrame(data = principalComponents
    , columns = ['principal component 1', 'principal component
↳ 2', 'principal component 3', 'principal component 4', 'principal component 5'])
finalDf=principalDf
```

```
[53]: #finalDf = pd.concat([principalDf, data[['fetal_health']]], axis = 1)
```

```
[54]: from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, confusion_matrix
from sklearn.metrics import classification_report
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
```

```
[55]: y.shape
```

```
[55]: (4965, 1)
```

```
[56]: finalDf.shape
```

```
[56]: (4965, 5)
```

```
[57]: X_train, X_test, y_train, y_test = train_test_split(finalDf, y, test_size=0.4)
```

```
[58]: X_train
```

```
[58]:      principal component 1  principal component 2  principal component 3  \
1587                0.956878                1.569415                0.588985
2343               -2.531951               -0.828686               -0.077501
2499               -1.368807                1.482100               -0.344704
4502                4.359033               -0.496193               -0.316583
1976               -1.492004                1.287881                1.256608
...
3276               -3.047368               -1.533019               -0.419620
4902                5.815971               -1.661513                0.439983
2641               -2.387667               -0.701829                0.287404
3016               -1.488183                0.712646               -0.675620
3876                3.328543               -1.426344                1.187249

      principal component 4  principal component 5
1587                1.228358                1.344599
2343               -0.286005               -0.253239
2499               -0.359640                0.394669
4502               -1.633983                1.208086
1976               -2.469014               -1.108613
...
3276               -0.270883               -0.137230
4902               -1.260138               -0.765890
2641               -0.104080               -0.544272
3016                1.389087                1.081168
3876                1.155247                1.347318
```

```
[2979 rows x 5 columns]
```

```
[84]: clf1 = RandomForestClassifier()
      clf1.fit(X_train, y_train)

      # Making predictions
      y1_pred = clf1.predict(X_test)

      # Evaluating the model
      acc1 = accuracy_score(y_test, y1_pred)
      print("Random forest Accuracy:", accuracy1)
      print(classification_report(y_test, y1_pred))
      conf_matrix = confusion_matrix(y_test, y1_pred)
```

```
plt.figure(figsize=(4, 3))
sns.heatmap(conf_matrix, annot=True, cmap='magma', fmt='g', cbar=False)

# Add labels, title, and adjust layout
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.title('Confusion Matrix')
plt.tight_layout()


# Show the plot
plt.show()
```

/tmp/ipykernel_33/3900602095.py:2: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
clf1.fit(X_train, y_train)
```

Random forest Accuracy: 0.9803625377643505

	precision	recall	f1-score	support
1.0	0.98	0.96	0.97	659
2.0	0.96	0.98	0.97	661
3.0	1.00	1.00	1.00	666
accuracy			0.98	1986
macro avg	0.98	0.98	0.98	1986
weighted avg	0.98	0.98	0.98	1986




```
[87]: # Create Decision Tree classifier object
      clf2 = DecisionTreeClassifier()

      # Train Decision Tree Classifier
      clf2 = clf2.fit(X_train,y_train)

      # Predict the response for test dataset
      y2_pred = clf2.predict(X_test)
      acc2=accuracy_score(y_test, y2_pred)
      # Model Accuracy, how often is the classifier correct?
      print("Decision tree Accuracy:",acc2 )
      print(classification_report(y_test, y2_pred))
```

Decision tree Accuracy: 0.9652567975830816

	precision	recall	f1-score	support
1.0	0.97	0.94	0.96	659
2.0	0.94	0.96	0.95	661
3.0	0.99	0.99	0.99	666
accuracy			0.97	1986
macro avg	0.97	0.97	0.97	1986
weighted avg	0.97	0.97	0.97	1986

```
[88]: # Create KNN classifier object
      clf3 = KNeighborsClassifier(n_neighbors=3) # You can adjust the number of
      ↪neighbors (k) here

      # Train KNN classifier
      clf3.fit(X_train, y_train)

      # Predict the response for test dataset
      y3_pred = clf3.predict(X_test)
      acc3=accuracy_score(y_test, y3_pred)
      # Model Accuracy, how often is the classifier correct?
      print("KNN Accuracy:", acc3)
      print(classification_report(y_test, y3_pred))
```

KNN Accuracy: 0.9813695871097684

	precision	recall	f1-score	support
1.0	1.00	0.95	0.97	659
2.0	0.95	0.99	0.97	661
3.0	1.00	1.00	1.00	666
accuracy			0.98	1986

macro avg	0.98	0.98	0.98	1986
weighted avg	0.98	0.98	0.98	1986

```
/opt/conda/lib/python3.10/site-
packages/sklearn/neighbors/_classification.py:215: DataConversionWarning: A
column-vector y was passed when a 1d array was expected. Please change the shape
of y to (n_samples,), for example using ravel().
    return self._fit(X, y)
```

```
[89]: # Create SVM classifier object
      clf4 = SVC()

      # Train SVM classifier
      clf4.fit(X_train, y_train)

      # Predict the response for test dataset
      y4_pred = clf4.predict(X_test)
      acc4=accuracy_score(y_test, y4_pred)
      # Model Accuracy, how often is the classifier correct?
      print("SVC Accuracy:",acc4 )
      print(classification_report(y_test, y4_pred))
```

```
SVC Accuracy: 0.9602215508559919
```

	precision	recall	f1-score	support
1.0	0.98	0.93	0.95	659
2.0	0.92	0.96	0.94	661
3.0	0.98	0.99	0.98	666
accuracy			0.96	1986
macro avg	0.96	0.96	0.96	1986
weighted avg	0.96	0.96	0.96	1986

```
/opt/conda/lib/python3.10/site-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
    y = column_or_1d(y, warn=True)
```

```
[90]: from sklearn.linear_model import LogisticRegression

      # Assuming X_train_pca and y_train are the PCA-transformed training data and
      ↳corresponding labels
      # Initialize logistic regression model
      lr = LogisticRegression()
```

```

# Train logistic regression on the transformed data
lr.fit(X_train, y_train)
y5_pred = lr.predict(X_test)
acc5=accuracy_score(y_test, y5_pred)
# Model Accuracy, how often is the classifier correct?
print("Logistic regression Accuracy:",acc5 )
print(classification_report(y_test, y5_pred))

```

Logistic regression Accuracy: 0.9330312185297079

	precision	recall	f1-score	support
1.0	0.95	0.93	0.94	659
2.0	0.90	0.90	0.90	661
3.0	0.95	0.96	0.95	666
accuracy			0.93	1986
macro avg	0.93	0.93	0.93	1986
weighted avg	0.93	0.93	0.93	1986

```

/opt/conda/lib/python3.10/site-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
y = column_or_1d(y, warn=True)

```

```

[91]: import matplotlib.pyplot as plt
import numpy as np

# Data
models = ['Random Forest', 'Decision Tree', 'Logistic Regression', 'SVC', 'KNN'] # Names of the models
accuracies = [acc1*100, acc2*100, acc3*100, acc4*100, acc5*100] # Accuracies of the models

# Create a figure and axis object
fig, ax = plt.subplots(figsize=(10, 6))

# Define width of each bar
bar_width = 0.35

# Index for the x-axis (position of each group)
index = np.arange(len(models))

# Plot the bars for accuracies
bars = ax.bar(index, accuracies, bar_width, label='Accuracy', color='skyblue')

```

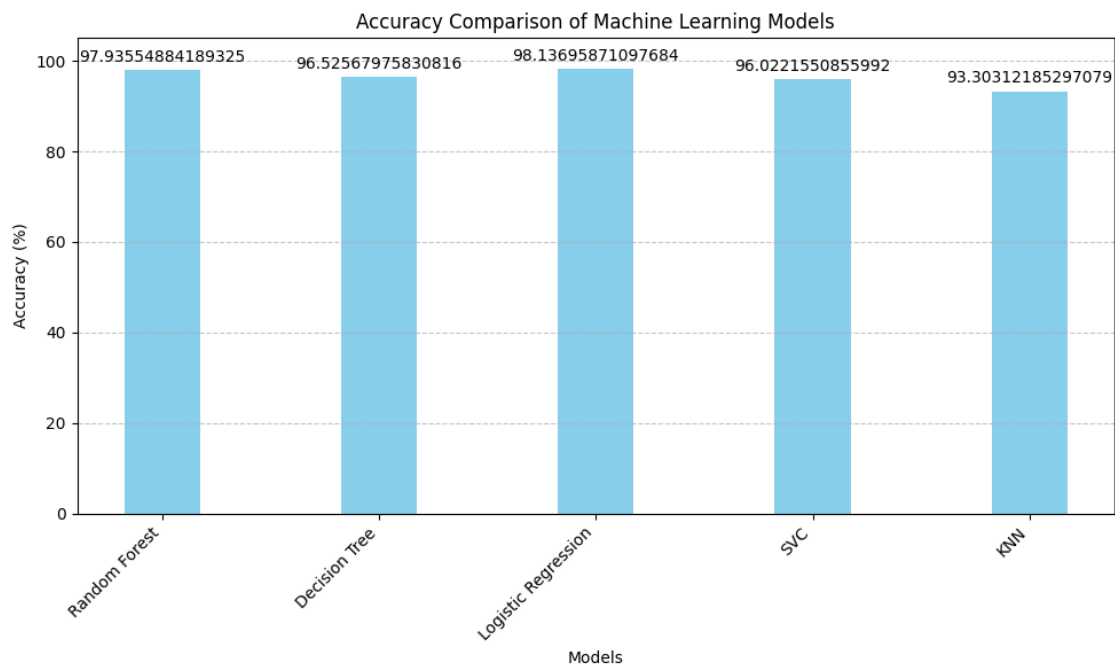
```

# Add labels, title, and grid
ax.set_xlabel('Models')
ax.set_ylabel('Accuracy (%)')
ax.set_title('Accuracy Comparison of Machine Learning Models')
ax.set_xticks(index)
ax.set_xticklabels(models, rotation=45, ha='right')
ax.set_ylim(0, 105)
ax.grid(axis='y', linestyle='--', alpha=0.7)

# Add labels on top of the bars
for bar in bars:
    height = bar.get_height()
    ax.annotate('{}' .format(height),
                xy=(bar.get_x() + bar.get_width() / 2, height),
                xytext=(0, 3), # 3 points vertical offset
                textcoords="offset points",
                ha='center', va='bottom')

plt.tight_layout()
plt.show()

```



```

[65]: p=data
      data
      p_x=data.drop('fetal_health',axis=1)
      p_y=data['fetal_health']

```

```
[66]: X_t, X_te, y_t, y_te = train_test_split(p_x, p_y, test_size=0.4)
      c1 = LogisticRegression()
      c1.fit(X_t, y_t)

      # Making predictions
      y_pr = c1.predict(X_te)

      # Evaluating the model
      acc1 = c1.score(X_te, y_te)
      print("Random forest Accuracy:", acc1)
      print(classification_report(y_te, y_pr))
```

```
Random forest Accuracy: 0.8719153936545241
              precision    recall  f1-score   support

    1.0         0.90        0.97        0.93        670
    2.0         0.60        0.39        0.47        111
    3.0         0.85        0.74        0.79         70

 accuracy                   0.87         851
 macro avg              0.78        0.70        0.73         851
 weighted avg           0.86        0.87        0.86         851
```

```
/opt/conda/lib/python3.10/site-packages/sklearn/linear_model/_logistic.py:458:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

```
Increase the number of iterations (max_iter) or scale the data as shown in:
    https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
    https://scikit-learn.org/stable/modules/linear_model.html#logistic-
regression
    n_iter_i = _check_optimize_result(
```